



By: Manish

Introduction to AGENT

An LLM agent is an AI system that uses a large language model (LLM) to perform tasks like reasoning, planning, and decision-making.





By: Manish

Prompt Engineering

It is the process of writing instructions to get the best results from an LLM models.

Prompting Techniques

- 1.Zero-Shot prompting
- 2.Few-Shot prompting
- 3.Chain-of-Thought (CoT)
- 4.Role Prompting


Prompt Engineering


1. Zero-shot Prompting

Prompting without any examples or demonstrations

Example Prompt:

Tell a joke about {animal}.





By: Manish



Prompt Engineering

2. Few-shot Prompting

Prompting with some examples or demonstrations

Example Prompt:

Here are some animal jokes to learn from:


1. Why don't dogs make good comedians?

Their tales are too "ruff"!

2. What do you call a bird that's afraid of heights?

A chicken!

Using these examples as inspiration, create a similar-style joke about {animal}.



By: Manish

Prompt Engineering

3. CoT Prompting

Step by step instructions and breaking down the writing process .

Example Prompt:

Let's create a joke about {animal}

Think step by step:

1. First, identify key characteristics of given animal
2. Let's find wordplay opportunities with given animal or related terms.
3. Create a setup for the joke
4. Craft the punchline using our wordplay

Please follow these steps to construct a joke and response final answer.

By: Manish

Prompt Engineering

4. Role Prompting

Providing the specific role and instructions before asking the questions, encouraging to get output from this role's viewpoint.

Example Prompt:

You are a professional comedian who specializes in animal-themed comedy. You've been performing at comedy clubs for 15 years and have mastered the art of animal jokes. You understand animals behaviour deeply and know how to craft jokes that resonate with both animal owners and general audiences.

Your style is:

- Clean humour (family-friendly)
- Clever wordplay
- Observational comedy about typical animal behaviours
- Timing and setup are crucial elements of your jokes

As an experienced animal comedy specialist, please create a joke about {animal}. Include your thought process about why this joke works and how it plays on common animal behaviours or characteristics.

Best Practices in Prompting

1.Be Specific

2.Use delimiters

3.Iterate: Test and refine prompt

2. Use delimiters like :

<Instructions>

Joke should be full of punchlines.

</Instructions>

<steps>

Analyze the topic of the joke

1.Don't ask like "Tell me a joke". Instead add some more detail like : "Tell me a joke about current issue of road constructions"

Find the relevant context for the joke

Create an setup


Add punchlines

</step>

Joke should be in below tag of response.

<response>...</response>

By: Manish



By: Manish

Prompt Engineering

Prompt engineering is the heart of the LLM models
Better the prompt better the outcome

References for prompt engineering

<https://www.deeplearning.ai/short-courses/chatgpt-prompt-engineering-for-developers/>

By: Manish

Making LLMs More Deterministic

Causes of non-deterministic responses:

LLMs generate responses through probabilistic sampling, where each word (token) is selected based on calculated probability distributions rather than fixed rules.

Ways to make more deterministic

By passing low temperature in LLM requests

By using more detailed instructed prompt

Why Non-deterministic is a problem

Can be problematic for consistent business applications.

Makes testing and validation more challenging.

May produce inconsistent result in critical applications.

By: Manish

Structured Output

Structured Response is the most needed feature for every developer right now .

In Early stage , developer have to write specific prompt to get desired response.


To overcome this type of unstructured response format , many frameworks have provided the structured output feature.

LLM Agent Workflow

LLM Agents combine language models with planning, memory, and tools to perform complex tasks

Fig: LLM Agent Workflow

By: Manish



By: Manish


Tools

Tools are the functions that are designed to perform certain task.

Examples of tools : Web Search, Calculator , API calls, File Operation

When to use??

When the agent needs to perform specific tasks beyond text generations.



Planning

Planning is used for reasoning and breaking down complex tasks into smaller,

manageable steps.

a. Prompt Templates

Structure the input to ensure consistent agent behavior.

Help maintain consistent formatting and instruction patterns. b.

Chain of Thought

Breaks down complex tasks into simpler steps.

Supports multi-step reasoning and complex problem-solving.

By: Manish

Memory

Short term memory

It maintains the conversation history for the execution of time and for n number of conversations.

Think of it like a whiteboard.

Long-term memory

Its like a diary ,storing the insights and information from the past interactions .

Good to use the vector database for this type of memory.

By: Manish

Knowledge Base

A knowledge base is structured or semi structured information that an agent can search, to improve its responses.

The knowledge base can take different forms, such as:

Static documents : PDFs,FAQs,...etc

Vector databases: Embedding based retrieval system

Databases : SQL or NoSQL databases

Vector Databases

Vector database offers the best solution for retrieving the relevant results from dense information quickly.

Steps Involved in AI Agent


1. Chunk the information
2. Generating Embedding
and Storing to Vector db
3. Query
4. Search the vector
database
5. Get meaningful output



Database used for embedding storage


ChromaDB
PgVector
MongoDB
LanceDB
Cassandra
Clickhouse
and many more

By: Manish



Chunking Approaches

Fixed Size Chunking
Semantic Chunking
Recursive Chunking



Agentic Chunking

Searching

keyword search


vector search

hybrid search

By: Manish

Working of Retrieval-Augmented Generation(RAG)





By: Manish

Working of RAG

Working Mechanism:

- 1.Retrieval : searches and retrieves the relevant documents or data based on the user query.
- 2.Augmentation: Retrieved information is combined with the user query to provide context.
- 3.Generation : LLM model uses this augmented input to generate accurate response.



By: Manish

Why use RAG ?

Benefits Of RAG


- 1.Prevent models from hallucinating
- 2.Allows model to cite sources
- 3.Expands the model use cases

4. Can stick to the internal knowledge base as required.

By: Manish

Use of RAG in AI Applications





By: Manish

LLM Agent Frameworks

Some of the popular frameworks:

Langchain's LangGraph

Agno

Autogen

CrewAi

Why use these??

Easy to switch LLM models

Can easily add knowledge base

Built-in Memory support

Custom tool support, along with some built-in tools



By: Manish

Let's Build Agentic AI Apps with Langchain

What we will do :

Set Up Environment with Langchain libraries and API keys

Create API Tools

Initialize Agent


Test the agent with and without tool call

By: Manish

Single agent System

A system where one autonomous entity manages the entire workflow from perception to action.

Acts as a "central brain" that reasons, plans, and executes tasks independently.



**SINGLE AGENT:
FOCUSED EXECUTION**



- AUTONOMOUSLY ACHIEVES A SPECIFIC GOAL (e.g., BOOK A FLIGHT)
- HIGH EFFICIENCY FOR DEFINED TASKS
- SELF-CONTAINED LEARNING & DECISION-MAKING


By: Manish

Multi-agent System

Combination of the agents to form a team

Imagine office hierarchy : under a project there will be multiple employee . And combination of them is Team.

Same in this too ,Multiple agents with different roles are combined to form an team to complete certain task.



Guardrails

A programmable layer that sits between the User and the AI Agent to monitor inputs and outputs.

Acts as a "filter" to ensure the AI stays within legal, ethical, and technical




boundaries.

By: Manish

Human-in-the-loop

A model where human intervention is integrated into the AI's decision-making process.

Humans act as a checkpoint to approve, reject, or refine the agent's actions.




By: Manish

Model context protocol (MCP)

An open standard that allows AI models to connect seamlessly to external data and tools.





By: Manish

Task

Open task : create any agentic system which contains tools use, vector database , multi agents.

By: Manish

References

<https://www.langchain.com/langgraph>

<https://developers.openai.com/api/docs/guides/agents>

<https://openai.github.io/openai-agents-python/quickstart/>

<https://docs.agno.com/introduction>

<https://learn.deeplearning.ai/courses/ai-agents-in-langgraph/lesson/qyrpc/introduction?courseName=ai-agents-in-langgraph>



Thank You

By: Manish

