



# Distributed Databases

Database System Concepts, 6<sup>th</sup> Ed.

©Silberschatz, Korth and Sudarshan

See [www.db-book.com](http://www.db-book.com) for conditions on re-use



# Distributed Database System

- A distributed database system consists of loosely coupled sites that share no physical component
- Database systems that run on each site are independent of each other
- Transactions may access data at one or more sites



# Homogeneous Distributed Databases

- In a homogeneous distributed database
  - All sites have identical software
  - Are aware of each other and agree to cooperate in processing user requests.
  - Each site surrenders part of its autonomy in terms of right to change schemas or software
  - Appears to user as a single system
- In a heterogeneous distributed database
  - Different sites may use different schemas and software
    - ▶ Difference in schema is a major problem for query processing
    - ▶ Difference in software is a major problem for transaction processing
  - Sites may not be aware of each other and may provide only limited facilities for cooperation in transaction processing



# Distributed Data Storage

- Assume relational data model
- Replication
  - System maintains multiple copies of data, stored in different sites, for faster retrieval and fault tolerance.
- Fragmentation
  - Relation is partitioned into several fragments stored in distinct sites
- Replication and fragmentation can be combined
  - Relation is partitioned into several fragments: system maintains several identical replicas of each such fragment.



# Data Replication

- A relation or fragment of a relation is **replicated** if it is stored redundantly in two or more sites.
- **Full replication** of a relation is the case where the relation is stored at all sites.
- Fully redundant databases are those in which every site contains a copy of the entire database.



# Data Replication (Cont.)

## ■ Advantages of Replication

- **Availability:** failure of site containing relation  $r$  does not result in unavailability of  $r$  if replicas exist.
- **Parallelism:** queries on  $r$  may be processed by several nodes in parallel.
- **Reduced data transfer:** relation  $r$  is available locally at each site containing a replica of  $r$ .

## ■ Disadvantages of Replication

- Increased cost of updates: each replica of relation  $r$  must be updated.
- Increased complexity of concurrency control: concurrent updates to distinct replicas may lead to inconsistent data unless special concurrency control mechanisms are implemented.
  - ▶ One solution: choose one copy as **primary copy** and apply concurrency control operations on primary copy



# Data Fragmentation

- Division of relation  $r$  into fragments  $r_1, r_2, \dots, r_n$  which contain sufficient information to reconstruct relation  $r$ .
- **Horizontal fragmentation**: each tuple of  $r$  is assigned to one or more fragments
- **Vertical fragmentation**: the schema for relation  $r$  is split into several smaller schemas
  - All schemas must contain a common candidate key (or superkey) to ensure lossless join property.
  - A special attribute, the tuple-id attribute may be added to each schema to serve as a candidate key.



# Horizontal Fragmentation of *account* Relation

| <i>branch_name</i> | <i>account_number</i> | <i>balance</i> |
|--------------------|-----------------------|----------------|
| Hillside           | A-305                 | 500            |
| Hillside           | A-226                 | 336            |
| Hillside           | A-155                 | 62             |

$$account_1 = \sigma_{branch\_name="Hillside"}(account)$$

| <i>branch_name</i> | <i>account_number</i> | <i>balance</i> |
|--------------------|-----------------------|----------------|
| Valleyview         | A-177                 | 205            |
| Valleyview         | A-402                 | 10000          |
| Valleyview         | A-408                 | 1123           |
| Valleyview         | A-639                 | 750            |

$$account_2 = \sigma_{branch\_name="Valleyview"}(account)$$



# Vertical Fragmentation of *employee\_info* Relation

| <i>branch_name</i> | <i>customer_name</i> | <i>tuple_id</i> |
|--------------------|----------------------|-----------------|
| Hillside           | Lowman               | 1               |
| Hillside           | Camp                 | 2               |
| Valleyview         | Camp                 | 3               |
| Valleyview         | Kahn                 | 4               |
| Hillside           | Kahn                 | 5               |
| Valleyview         | Kahn                 | 6               |
| Valleyview         | Green                | 7               |

$\text{deposit}_1 = \Pi_{\text{branch\_name}, \text{customer\_name}, \text{tuple\_id}}(\text{employee\_info})$

| <i>account_number</i> | <i>balance</i> | <i>tuple_id</i> |
|-----------------------|----------------|-----------------|
| A-305                 | 500            | 1               |
| A-226                 | 336            | 2               |
| A-177                 | 205            | 3               |
| A-402                 | 10000          | 4               |
| A-155                 | 62             | 5               |
| A-408                 | 1123           | 6               |
| A-639                 | 750            | 7               |

$\text{deposit}_2 = \Pi_{\text{account\_number}, \text{balance}, \text{tuple\_id}}(\text{employee\_info})$



# Advantages of Fragmentation

- Horizontal:
  - allows parallel processing on fragments of a relation
  - allows a relation to be split so that tuples are located where they are most frequently accessed
- Vertical:
  - allows tuples to be split so that each part of the tuple is stored where it is most frequently accessed
  - tuple-id attribute allows efficient joining of vertical fragments
  - allows parallel processing on a relation
- Vertical and horizontal fragmentation can be mixed.
  - Fragments may be successively fragmented to an arbitrary depth.



# Query Transformation

- Translating algebraic queries on fragments.
  - It must be possible to construct relation  $r$  from its fragments
  - Replace relation  $r$  by the expression to construct relation  $r$  from its fragments
- Consider the horizontal fragmentation of the *account* relation into
  - $account_1 = \sigma_{branch\_name = "Hillside"}(account)$
  - $account_2 = \sigma_{branch\_name = "Valleyview"}(account)$
- The query  $\sigma_{branch\_name = "Hillside"}(account)$  becomes
$$\sigma_{branch\_name = "Hillside"}(account_1 \cup account_2)$$
which is optimized into
$$\sigma_{branch\_name = "Hillside"}(account_1) \cup \sigma_{branch\_name = "Hillside"}(account_2)$$



## Example Query (Cont.)

- Since  $account_1$  has only tuples pertaining to the Hillside branch, we can eliminate the selection operation.
- Apply the definition of  $account_2$  to obtain
$$\sigma_{branch\_name = "Hillside"} (\sigma_{branch\_name = "Valleyview"} (account))$$
- This expression is the empty set regardless of the contents of the *account* relation.
- Final strategy is for the Hillside site to return  $account_1$  as the result of the query.



# Heterogeneous Distributed Databases

- Many database applications require data from a variety of preexisting databases located in a heterogeneous collection of hardware and software platforms
- Data models may differ (hierarchical, relational, etc.)
- Transaction commit protocols may be incompatible
- Concurrency control may be based on different techniques (locking, timestamping, etc.)
- System-level details almost certainly are totally incompatible.
- A **multidatabase system** is a software layer on top of existing database systems, which is designed to manipulate information in heterogeneous databases
  - Creates an illusion of logical database integration without any physical database integration



# Query Processing

- Several issues in query processing in a heterogeneous database
- Schema translation
  - Write a **wrapper** for each data source to translate data to a global schema
  - Wrappers must also translate updates on global schema to updates on local schema
- Removal of duplicate information when sites have overlapping information
  - Decide which sites to execute query



# Mediator Systems

- **Mediator** systems are systems that integrate multiple heterogeneous data sources by providing an integrated global view, and providing query facilities on global view
  - Unlike full fledged multidatabase systems, mediators generally do not bother about transaction processing
  - But the terms mediator and multidatabase are sometimes used interchangeably
  - The term **virtual database** is also used to refer to mediator/multidatabase systems



# Directory Systems

- Typical kinds of directory information
  - Employee information such as name, id, email, phone, office addr, ..
  - Even personal information to be accessed from multiple places
    - ▶ e.g., Web browser bookmarks
- White pages
  - Entries organized by name or identifier
    - ▶ Meant for forward lookup to find more about an entry
- Yellow pages
  - Entries organized by properties
  - For reverse lookup to find entries matching specific requirements



# Directory Access Protocols

- Most commonly used directory access protocol:
  - LDAP (Lightweight Directory Access Protocol)
  - Simplified from earlier X.500 protocol
- Question: Why not use database protocols like ODBC/JDBC?
- Answer:
  - Simplified protocols for a limited type of data access, evolved parallel to ODBC/JDBC
  - Provide a nice hierarchical naming mechanism similar to file system directories
    - ▶ Data can be partitioned amongst multiple servers for different parts of the hierarchy, yet give a single view to user
      - E.g., different servers for Bell Labs Murray Hill and Bell Labs Bangalore
    - Directories may use databases as storage mechanism



# **End of Course**

**Database System Concepts, 6<sup>th</sup> Ed.**

**©Silberschatz, Korth and Sudarshan**

**See [www.db-book.com](http://www.db-book.com) for conditions on re-use**